



1.2cm]

Complexity-Aware Energy Estimation

6pt] *for an LLM Data Curation Pipeline* 1cm]

Master Thesis

6pt] *Master of Science in Applied Data Science* 4pt] *Utrecht University* 1cm]

Author:

4pt] *Romir Malik* 4pt] `r.malik@students.uu.nl` 1cm]

Supervision:

4pt] *Prof. Dr. Martino Mensio* (*Utrecht University, first supervisor*) 4pt] *Dr. Lisa By*

Date:

4pt] *July 2026* 0.8cm]

0.6cm]

Abstract. Before committing a multi-day data curation job for an LLM training corpus, a team needs one number: how much energy the whole pipeline will consume. This thesis builds that predictor for the GPT-NL curation pipeline on the Snellius supercomputer across four diverse corpora spanning a 100:1 range in document length (1,798 to 181,075 characters per document).

The methodology is a sample-then-predict framework. Run a few small calibration sizes on exclusive nodes (minutes, not hours). Fit a two-parameter linear energy model per stage. Predict the expensive full run with a closed-form prediction interval. A sequential Bayesian estimator walks the live pipeline and refines the estimate: clean telemetry readings tighten the intervals, contaminated or missing readings are rejected by an automatic innovation gate. A cross-corpus transfer model g predicts the energy coefficients for a corpus that was never measured, using features derived from the data itself.

Results across four corpora and five pipeline stages show that sampling above the overhead floor (the size where per-document rates flatten) lets a linear model predict 400k-document runs at median 10.3% error for the dominant stage. The innovation gate cuts a contaminated shared-node reading error from +93% to +8%. Energy per character is stable at approximately 2.2×10^{-4} J/char across the string normalization stage, making the predictor portable. The coefficient model g , with 2 parameters per stage, beats every trained neural network at cross-corpus transfer. The total pipeline at 400k documents consumes 2.2 MJ for short-document corpora and scales with both document count and document length.

Declaration on the Use of Generative AI

In accordance with Utrecht University guidelines, I declare that generative AI tools were used in the preparation of this thesis. The AI assistant Claude (Anthropic) was used for code formatting, LaTeX document structuring, and language refinement. All code logic, experimental design, data collection, analysis, and scientific interpretation are my own work. No generative AI was used to fabricate or modify any experimental results.

2pt]Romir Malik, July 2026

Contents

Declaration on the Use of Generative AI	1
1 Introduction	3
2 Data	5
3 Method: The Sample-Predict Framework	6
3.1 Per-stage energy model	6
3.2 Where to sample: the overhead floor	7
3.3 The stopping rule	8
3.4 The sequential estimator	9
3.5 The innovation gate	9
3.6 Time-driven estimation on shared nodes	10
3.7 Cross-corpus transfer: the coefficient model g . .	10
4 Results	11
4.1 Per-corpus calibration	11
4.2 Sample-then-predict: held-out results	11
4.3 Pipeline total across corpora	12
4.4 The innovation gate	12
4.5 Shared node contamination	12
4.6 Time-driven estimation	13
4.7 Cross-corpus transfer	13
4.8 The coefficient model g	14
4.9 Deduplication	14
4.10 Model comparison	14
4.11 Live monitor	15
4.12 Forecast tool	16
5 Limitations	16
6 Conclusion	16

1 Introduction

Energy reporting for large language models concentrates on training and inference [8, 7]. The step that comes before both—data curation—is run over hundreds of millions of documents to filter, normalise, deduplicate, and clean the training corpus. Its energy cost is rarely reported and to our knowledge never predicted before a run.

This thesis was conducted within the GPT-NL project, a Dutch-English foundation model developed from the ground up for responsible deployment in critical sectors [10]. Unlike models trained indiscriminately on internet-scale data, GPT-NL prioritises verifiable sources, data privacy, and compliance with the EU AI Act and GDPR. The project’s data pipeline has two major phases: data extraction from raw sources (out of scope for this work), and data curation—a multi-stage processing chain that transforms extracted documents into a training-ready corpus.

The curation pipeline runs on the Snellius national supercomputer [9] using the Datatrove framework [4]. It consists of six stages in production (data splitting, string normalisation, heuristic filtering, PII masking, deduplication, toxic language detection), of which we measure five. The cluster is instrumented with the Energy Aware Runtime (EAR), a hardware-level monitoring system developed at the Barcelona Supercomputing Center [1]. EAR records per-job metrics through SLURM integration: DC node power, CPU utilisation, memory bandwidth, and instructions per cycle. On Snellius’s AMD EPYC (Genoa) nodes, DRAM power is not separately instrumented—only total DC node power is available.

The practical need is concrete. Before launching a multi-day curation job over a new corpus, the team wants a whole-pipeline energy estimate with an honest error bar, obtained from small trial runs that finish in minutes. Two problems make this hard.

First, node-level DC power on a shared HPC cluster includes

co-tenant workloads. A controlled experiment showed the same stage with identical work, identical wall time, and identical CPU utilisation but 2.4 times higher measured power when run on a shared node compared to an exclusive node. Second, the deduplication stage runs as four chained sub-jobs that produced no EAR records in the initial cluster configuration: the sub-job launcher never created the directory that EAR writes its per-job database into, so the writes silently failed. This was a single-line fix—adding a `mkdir` call before the sub-job launcher—highlighting how fragile energy telemetry infrastructure can be in practice.

We calibrate a two-parameter linear energy model per stage from small exclusive-node runs, then walk the pipeline with a sequential Bayesian estimator whose innovation gate decides whether to trust each incoming reading or fall back to the model. The main contributions are:

1. A full pipeline energy account for all five stages across four corpora spanning a 100x range in document length. Sampling above the overhead floor is identified as the condition that makes calibration reliable.
2. A sample-then-predict framework: the closed-form prediction interval from a few small runs produces held-out error of 10.3% for the dominant stage at 400k documents. A recursive stopping rule tells the user when enough samples have been taken.
3. A telemetry-robust sequential estimator. The innovation gate cuts contaminated reading error from +93% to +8%.
4. A time-driven estimator ($E = P_{\text{eff}} \times t$) that resolves the shared-node contamination problem, reducing error from 45.8% to 5.4%.
5. A coefficient transfer model g that predicts energy coefficients for unseen corpora from data-derived features and beats every trained neural network at cross-corpus transfer.
6. Two practical tools: a forecast tool (energy, cost, CO2 in one command) and a live monitoring dashboard.

2 Data

All measurements were collected on the Snellius genoa partition (AMD EPYC 9654, 192 cores per node). We ran the GPT-NL curation pipeline [10] implemented with Datatrove [4] across four corpora from the GPT-NL Public Corpus [3]:

- American Stories [2]: historical US newspaper text, 1,798 characters per document. 8 sizes from 1k to 1.4M documents.
- GitHub open source code: 3,151 characters per document. 7 sizes from 1k to 700k documents.
- German public domain: 48,881 characters per document. 7 sizes from 1k to 200k documents.
- European Parliament [5]: 181,075 characters per document. 9 sizes from 100 to 400k documents.

These four corpora span a 100:1 range in document length. Characters per document were measured by dividing the total I/O bytes read during data splitting by the number of output documents — a data-derived measurement that does not require probing the corpus separately.

Each size was run with multiple replicates at single-task (ntasks=1) and multi-task (ntasks=16, 32) configurations. The production curation pipeline [10] consists of six stages, of which we measured five (PII masking requires the PrivateAI service and was excluded from this study):

1. Data splitting: shards the raw corpus into per-document records. Lightweight, I/O bound.
2. String normalization: per-document Unicode normalisation (NFC). Memory intensive, character-level processing.
3. Heuristic filtering: FastText language detection and six quality filters (language, digit fraction, symbol ratio, repetition, document length, stop words). The most compute-intensive CPU stage.
4. Deduplication: MinHash near-duplicate removal with four chained sub-jobs (signature, buckets, cluster, filter).

5. Toxic language detection: GPU stage running on the H100 partition. Included for completeness.

The full dataset contains 442 aggregated measurements from 126 distinct runs across corpora, sizes, and parallelism levels.

EAR was configured with the monitoring policy (as opposed to minimization, which actively adjusts CPU frequency). The monitoring policy records energy metrics without interfering with job execution. EAR produces two output files per node per job: a per-job aggregate file (`_apps.csv`) containing total energy and average power, and a time-series file (`_loops.csv`) with per-loop metrics. We use the aggregate files exclusively, summing $\text{DC_NODE_POWER_W} \times \text{ELAPSED}$ across all rows to obtain per-stage energy in Joules. On AMD Genoa nodes, `DRAM_POWER_W` is not available—only total DC node power is recorded—so we work with node-level energy rather than component-level breakdowns.

3 Method: The Sample-Predict Framework

Figure 1a shows the full estimation pipeline as a flow chart. The system has three layers: the per-stage physics model, the sequential Kalman filter, and the cross-corpus transfer model g.


3.1 Per-stage energy model

Each stage s processes n input documents. Its energy follows a linear relationship:

$$E_s(n) = c_0^{(s)} + c_1^{(s)}n + \varepsilon_s \quad (1)$$

where c_0 is the fixed startup cost and c_1 is the marginal energy per document. We fit this from small exclusive runs using ridge regression on the slope only [6]. The prediction interval for a new size n_* is:

$$\hat{E}_s(n_*) \pm t_{\alpha/2} \cdot \hat{\sigma}_s \sqrt{1 + x_*^\top (X^\top X + \lambda I')^{-1} x_*} \quad (2)$$



../poster.png

(a) The poster summarising the full framework. The flow chart at the top shows how cheap trials feed into a swappable model slot per stage, tied together by the EKF. The side panel shows the mathematics. The bottom panels show the real findings across all four corpora.

This closed-form interval is exact and propagates cleanly through the pipeline, the reason we keep the model linear.

3.2 Where to sample: the overhead floor

A natural question is how many calibration runs are needed.

The more important point is where to place them. The per-document rate is:

$$r(n) = \frac{E_s(n)}{n} = c_1 + \frac{c_0}{n} \quad (3)$$

This rate falls toward the true marginal cost c_1 as n grows. Below a size called the overhead floor, $n^* \approx c_0/c_1$, the fixed cost

dominates and the slope is unidentifiable. Sampling below the floor produces a biased slope estimate and poor extrapolation. The rule is to sample above the overhead floor. For heuristic filtering on American Stories, this floor is approximately 40k documents.

3.3 The stopping rule

The stopping rule tells the user when enough calibration runs have been taken. After each new sample, the prediction interval for the full target run is recomputed. The relative half-width of this interval shrinks monotonically as more trials accumulate. We stop when:

$$\frac{1.96\sqrt{\text{Var}(\hat{E}(N))}}{\hat{E}(N)} \leq \tau \quad (4)$$

where τ is a tolerance. With $\tau = 0.10$ (10% relative interval width), the stop test fires after 3 samples on the dominant stage, with the predicted 400k-document run already within 7% of the true value.




Figure 2: The stopping rule in action. Each calibration trial tightens the 95% confidence ribbon. The stop test fires at 3 trials when the ribbon drops below 10% relative width, with the prediction already within 7% of the true full run.

3.4 The sequential estimator

The pipeline executes each stage as M_s parallel SLURM tasks. After task k of stage s completes, the EAR library records its energy z_k . We maintain a per-stage energy estimate $\hat{E}^{(s)}$ that blends the calibrated model prior with the streaming data using a Bayesian update rule.

The estimate after k of M_s tasks of stage s have reported is:

$$\hat{E}_k^{(s)} = w_k \cdot \hat{E}_{\text{extrap}}^{(s)} + (1 - w_k) \cdot \hat{E}_0^{(s)} \quad (5)$$

where $\hat{E}_{\text{extrap}}^{(s)} = (\sum_{i=1}^k z_i) \cdot M_s / k$ is the data-only extrapolation, $w_k = k / (k + k_0)$ is the blending weight, and $k_0 = 2$ controls how quickly the estimator trusts the arriving data over the model prior. When $k = 0$, the estimate is the pure model prior. When $k = M_s$, the estimate equals the measured sum with only instrument noise.

The variance of the estimate shrinks as more tasks report:

$$\text{Var}(\hat{E}_k^{(s)}) = (1 - w_k)^2 \cdot P_0^{(s)} + \sigma_{\text{sample}}^2 \cdot \max(M_s - k, 0) \quad (6)$$

with prior variance $P_0^{(s)} = (0.45 \cdot \hat{E}_0^{(s)})^2$ (45% relative uncertainty before any readings) and σ_{sample}^2 estimated from the variance of the k accepted readings. The total pipeline variance sums across stages, and the 95% confidence interval is

$$\pm 1.96 \sqrt{\sum_s \text{Var}(\hat{E}^{(s)})}.$$

3.5 The innovation gate

A per-task reading is accepted only if it is physically plausible—a single task cannot consume more energy than the entire stage’s predicted budget:

$$\frac{z_k}{\hat{E}_0^{(s)}} \leq \gamma, \quad \gamma = 2.5 \quad (7)$$

If the gate triggers, the contaminated reading is replaced by the model’s per-task share $\hat{E}_0^{(s)}/M_s$. The gate targets shared-node contamination where co-tenant workloads inflate a node’s power reading (up to $2.4\times$, Section 5) without affecting its wall-clock time. On clean exclusive-node runs, the gate is never observed to fire on the dominant stage.

3.6 Time-driven estimation on shared nodes

On exclusive nodes, the per-stage effective power $P_s^{\text{eff}} = E_s/t_s$ is stable at approximately 290 W on genoa. Wall time is co-tenancy-invariant while node power is not. We replace the contaminated power signal with:

$$\hat{E}_s = P_s^{\text{eff}} \cdot t_s \quad (8)$$

3.7 Cross-corpus transfer: the coefficient model g

Within one corpus, document count n and total character count $C = n \cdot \bar{\ell}$ are indistinguishable because $\bar{\ell}$ is constant. A second corpus with a different $\bar{\ell}$ breaks this collinearity.

However, the per-character coefficient $k = c_1/\bar{\ell}$ is only one feature. The model g generalises this by predicting the energy law’s slope c_1 for a new corpus from multiple features derived from the data itself:

$$\widehat{\log_{10} c_1^{(s)}} = g_s(\text{chars per doc, survival, cpi, io, gflops, util})$$

The label is the fitted slope $c_1^{(s)}$ per stage per corpus.

Characters per document come from the IO rate in data splitting. Survival is the keep rate of the quality filter. The rest are per-stage compute counters from the EAR telemetry.

Fitting in log space is necessary because coefficients span orders of magnitude across corpora.

The honest test is leave-one-corpus-out: train g on three corpora, predict the held-out corpus, apply the physics equation, and score the whole-run prediction.

4 Results

4.1 Per-corpus calibration

Table 1 shows the fitted coefficients for each corpus and stage.

The coefficients reveal how energy profiles differ across corpora. String normalization on American Stories costs 0.40 J/doc. On German public domain, with 48,881 characters per document, it costs 12.5 J/doc. The energy ratio ($31\times$) approximately tracks the document-length ratio ($27\times$), consistent with a per-character cost that is roughly stable.

Corpus	Stage	c_0 (J)	c_1 (J/doc)	R^2
amnews	Data splitting	4,975	0.029	0.950
amnews	String normalization	15,708	0.436	0.988
amnews	Heuristic filtering	138,452	3.627	0.978
amnews	Deduplication	20,053		0.791
github	Data splitting	807	0.049	0.972
github	String normalization	24,608		1.102
github	Heuristic filtering	-3,394	1.083	0.960
github	Deduplication	21,473		0.012
german	Data splitting	1,245	0.243	0.999
german	String normalization	155,157	25.071	0.978
german	Heuristic filtering	118,282	80.424	0.992
german	Deduplication	27,989	9.809	0.997
euparl	Data splitting	4,273	0.061	0.064
euparl	String normalization	169,085	8.112	0.385
euparl	Heuristic filtering	430,727	18.770	0.395
euparl	Deduplication	83,136	2.427	0.410

Table 1: Per-corpus and per-stage linear model fits from exclusive single-task runs. The coefficients span two orders of magnitude across corpora.

4.2 Sample-then-predict: held-out results

The real test is: calibrate on small sizes, predict a large run.

Table 2 shows this for 400k documents with calibration on sizes up to 100k.

Corpus	Stage	True (kJ)	Predicted (kJ)	MRE
amnews	Heuristic filtering	1,631	1,799	10.3%
amnews	String normalization	214	183	14.1%
amnews	Data splitting	19	12	35.0%
amnews	Deduplication	191	104	45.6%
github	Data splitting	17	17	0.6%
github	String normalization	655	335	48.8%
github	Heuristic filtering	548	249	54.6%

Table 2: Held-out prediction at 400k documents. Calibration on sizes up to 100k. The dominant stage (heuristic filtering on amnews) predicts at 10.3% MRE.

The results vary by corpus and stage. American Stories heuristic filtering predicts well at 10.3% MRE because its marginal cost is stable across sizes. GitHub string normalization misses because code text has different structure at scale. The overhead floor (Section 3.1) explains why: stages with large startup costs relative to their per-document rate require the calibration sizes to exceed that floor.

4.3 Pipeline total across corpora

Table 3 shows the complete pipeline energy at common operating points. The total varies dramatically by corpus due to document length.

Corpus	Size	Total (kJ)	Dominant stage	Share
amnews	100k	627	Heuristic filtering	75%
amnews	400k	2,169	Heuristic filtering	75%
github	100k	226	String normalization	36%
github	400k	1,343	String normalization	49%
german	100k	10,923	Heuristic filtering	71%
german	200k	4,892	String normalization	99%
euparl	40k	1,531	Heuristic filtering	64%

Table 3: Complete pipeline energy at common sizes. The dominant stage flips by corpus.

The dominant stage flips between corpora. For short documents (amnews, github), heuristic filtering dominates. For long documents (german), string normalization dominates because character-level operations scale with document length.

4.4 The innovation gate

Mode	Result
A (pre-run forecast)	Total 273 kJ, 95% CI [208, 338], zero telemetry
B (online, clean readings)	Converges to 253 kJ, CI width 65 to 7 kJ
Contaminated reading ingested	+93% error without gate
Gated EKF with contamination	+8% error

Table 4: EKF gate performance on a four-stage 100k-document pipeline walk.




Figure 3: The EKF walking the pipeline. Left: forecast and 95% band tighten as clean readings arrive. Right: a contaminated reading is rejected by the gate.

4.5 Shared node contamination

The controlled experiment confirmed the source of the problem.

Power (W)	Wall time (s)	CPU util	Energy (kJ)
Exclusive	275	595 97%	164
Shared	669	595 97%	398

Table 5: Controlled experiment: exclusive versus shared node. Wall time is identical, power is 2.4x higher on the shared node.

4.6 Time-driven estimation

Stage	P^{eff} (W)	Power-based MRE	Time-based MRE
Data splitting	293	54.2%	12.9%
String normalization	297	20.3%	3.0%
Heuristic filtering	278	63.1%	0.2%
Overall	290	45.8%	5.4%

Table 6: Shared node prediction error: power-based versus time-driven.




Figure 4: Left: stage energy versus wall time at $R^2 = 0.99$. Right: time-driven estimate cuts error from 45.8% to 5.4%.

4.7 Cross-corpus transfer

The per-character coefficients vary more than a simple constant model would predict. Within string normalization, k spans one order of magnitude (2.43×10^{-5} to 2.55×10^{-4} J/char).

The per-character transfer law provides a first-order approximation (mean k predicts within 60% MRE on held-out corpora, Section 8), but corpus-specific factors—character encoding density, Unicode normalisation complexity, and memory subsystem behaviour—cause real differences. The learned model g captures these by incorporating compute counters (CPI, GFLOPS, I/O throughput) as additional features.

Stage	amnews	github	german	euparl					
String normalization	2.24e-4	1.73e-4	2.55e-4	2.43e-5	Heuristic filtering	8.18e-4	1.71e-4		
	8.79e-4	5.16e-5	Deduplication	1.63e-4	1.79e-6	1.08e-4	5.14e-6		

Table 7: Per-character energy coefficient $k = c_1/\ell$ (J/char) across four corpora. Computed from ntasks=1 exclusive-node measurements.

4.8 The coefficient model g

The coefficient model g was evaluated using leave-one-corpus-out cross-validation. The comparison is against the per-character baseline (assuming k is constant at the mean of the training corpora).

Target	Baseline MRE	g MRE					
	Wall time	94.2%	85.3%	Energy	108.3%	87.9%	

Table 8: Leave-one-corpus-out cross-validation: g versus the per-character baseline. g wins on both targets.

The learned g beats the per-character baseline on both targets, proving the extra features genuinely help predict an unseen corpus. The absolute error is still high because leaving one corpus out leaves only three to train on.

4.9 Deduplication

Deduplication runs as four chained sub-jobs. None of the four produced an EAR record until a single mkdir command fixed the missing EAR directory. With telemetry restored, the signature pass dominates at scale.

Documents	Dedup energy (kJ)	Signature share							
1k	9.9	35%	100k	34.8	73%	400k	125.6	84%	

Table 9: Deduplication energy on clean exclusive nodes. The signature pass grows from 35% to 84% of stage energy as size increases.

At 1.4M documents the signature pass runs out of memory. This is the one place the pipeline breaks, a scaling limit that only appears when running the complete pipeline.

4.10 Model comparison

We compared OLS, Ridge, GBM, MLP, and FT-Transformer on two tasks: size extrapolation and corpus transfer.

Model	Size extrapolation MRE	Corpus transfer MRE
Linear (OLS, naive)	278% 483%	Ridge 277% 479% GBM 153% 330% MLP 111%
	3,272%	FT-Transformer 101% 362%
Per-stage physics + g (this work)		
	10.3%	87.9%

Table 10: Model comparison on energy prediction. Metric is mean relative error.

The naive OLS/GBM/MLP/FT-Transformer models are trained on a flat feature table without the per-stage physics structure. They perform poorly on both tasks because they must learn both the linear energy law and the per-stage identity from data. In contrast, the per-stage physics model (this work) applies the calibrated $E = c_0 + c_1 n$ law per stage with the coefficient model g for cross-corpus transfer, achieving 10.3% MRE on size extrapolation and 87.9% MRE on corpus transfer. The key architectural difference is that the physics backbone constrains the prediction to a known functional form, reducing the learning problem to coefficient prediction rather than function discovery.

4.11 Live monitor




Figure 5: Live monitor dashboard. The EKF estimate and 95% confidence band converge as each per-task reading arrives.

The live monitor (Figure 5) shows the EKF estimator running in real time. Each dot is a per-task energy reading. The band starts wide at the pure model prediction and collapses as clean readings arrive. Contaminated readings are flagged and rejected.

4.12 Forecast tool

The forecast tool translates predictions into operational numbers one command:

```
python3 forecast.py --n 400000
```

At 400k documents on American Stories, the tool reports: total pipeline 1.0 MJ (0.33 kWh with PUE 1.20), energy cost EUR 0.10 at EUR 0.30/kWh, carbon 0.10 kg CO₂, per-stage breakdown with 95% confidence bands, and a warning if deduplication exceeds its memory limit at 1.4M documents.

5 Limitations

Corpus count is the binding constraint. The twelve Dutch corpora were queued but emitted no usable energy telemetry. The transfer model g trains on four corpora. Section 10 shows that corpus count, not model size, limits cross-corpus transfer.

Replicate depth. Two to three replicates per size per stage gives noisy estimates. More replicates would tighten prediction intervals.

Deduplication memory. The 1.4M deduplication run fails out of memory in the signature pass. The linear model is validated up to 400k for this stage.

GPU stages. Toxic language detection on the H100 partition is included but with limited calibration sizes.

6 Conclusion

We built a whole-pipeline energy predictor for the GPT-NL data curation pipeline across four diverse corpora. The methodology is a sample-then-predict framework: run cheap calibration trials above the overhead floor, fit a per-stage linear model, predict the full run with a closed-form interval, then walk the live pipeline with a sequential Bayesian estimator that gates contaminated readings.

The key findings are:

The sequential estimator, not a full Extended Kalman Filter, blends model predictions with live telemetry using a Bayesian update rule parameterised by the number of completed tasks per stage. This design was chosen because per-stage energies are independent in the pipeline architecture (each stage runs to completion before the next begins), making a full state transition model unnecessary.

- The overhead floor (n^* approximately 40k for the dominant stage) is identified as the condition that makes the linear model reliable.
- Sampling above the floor lets a linear model predict 400k-document runs at 10.3% MRE for heuristic filtering.
- The stopping rule fires after 3 trials, with the prediction already within 7% of the true run.
- The innovation gate cuts contaminated reading error from +93% to +8%.
- Time-driven estimation ($E = P_{\text{eff}} t$) reduces shared-node error from 45.8% to 5.4%.
- Energy per character varies from 2.4×10^{-5} to 8.8×10^{-4} J/char across corpora and stages, with string normalization showing the tightest spread (one order of magnitude).
- The coefficient model g beats every trained neural network at cross-corpus transfer.
- The total pipeline varies by corpus: 2.2 MJ for American Stories at 400k documents, scaling with document count and document length across corpora.
- The dominant stage flips: heuristic filtering for short documents, string normalization for long documents.
- Deduplication required a telemetry fix and hits a memory limit at 1.4M documents.
- Two practical tools (forecast and live monitor) make the method usable in production.

The bottleneck for cross-corpus transfer is corpus count, not model size. The twelve Dutch corpora that did not emit

usable telemetry are the priority expansion path.

References

- [1] Barcelona Supercomputing Center. Energy aware runtime (EAR). https://gitlab.bsc.es/ear_team/ear, 2024. Accessed: 2026-06-09.
- [2] Melissa Dell, Jacob Carlson, Tom Bryan, Emily Silcock, Abhishek Arora, et al. American Stories: A large-scale structured text dataset of historical U.S. newspapers. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. Datasets and Benchmarks Track.
- [3] GPT-NL Consortium. GPT-NL public corpus. https://huggingface.co/datasets/GPT-NL/GPT-NL_Public_Corpus, 2025. 29 curated collections, 524B tokens, permissively licensed.
- [4] HuggingFace. Datatrove: Large scale data processing for LLM training. <https://github.com/huggingface/datatrove>, 2024. Accessed: 2026-06-09.
- [5] Philipp Koehn. Europarl: A parallel corpus for statistical machine translation. In *MT Summit X*, pages 79–86, 2005. URL <https://www.statmt.org/europarl/>.
- [6] Douglas C Montgomery, Elizabeth A Peck, and G Geoffrey Vining. *Introduction to Linear Regression Analysis*. John Wiley & Sons, 5th edition, 2012.
- [7] David Patterson, Joseph Gonzalez, Quoc Le, Chen Liang, Lluís-Miquel Munguia, Daniel Rothchild, David So, Maud Texier, and Jeff Dean. Carbon emissions and large neural network training. 2021. doi: 10.48550/arXiv.2104.10350. URL <https://arxiv.org/abs/2104.10350>.
- [8] Emma Strubell, Ananya Ganesh, and Andrew McCallum. Energy and policy considerations for deep learning in NLP. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 3645–3650, 2019. doi: 10.18653/v1/P19-1355. URL <https://aclanthology.org/P19-1355>.

- [9] SURF. Snellius national supercomputer. <https://www.surf.nl/en/snellius-national-supercomputer>, 2024. Accessed: 2026-06-09.
- [10] TNO and SURF. GPT-NL data curation pipeline architecture. Technical Report GPTNL-DEL-4001, TNO, 2025. TNO Public.